

P8106 HOMEWORK 3

xc2474 Xinlei Chen

4/9/2019

Problem

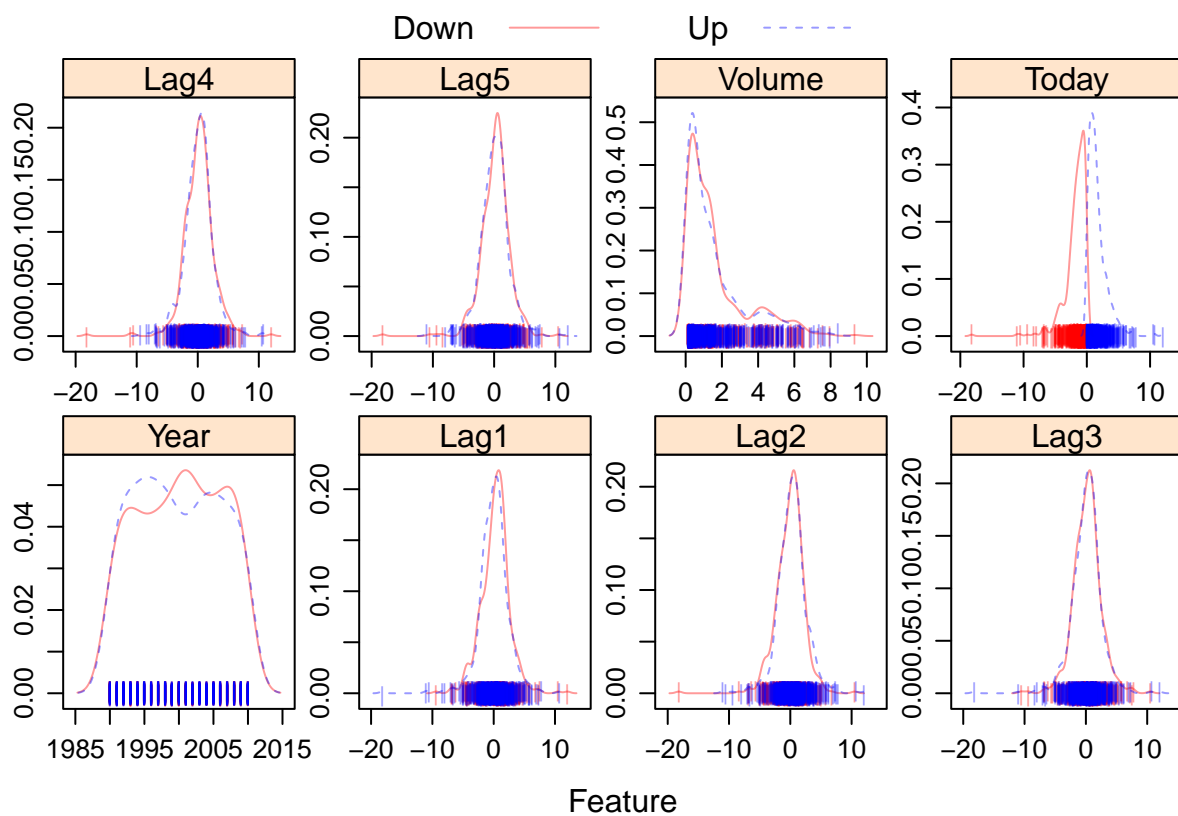
This questions will be answered using the *Weekly* data set, which is part of the *ISLR* package. This data is similar in nature to the *Smarket* data on the textbook except that it contains 1,089 weekly returns for 21 years, from the beginning of 1990 to the end of 2010. A description of the data can be found by typing `?Weekly` in the Console. (Note that the column *Today* is not a predictor.)

```
# load packages
library(tidyverse)
library(ISLR)
library(caret)
library(AppliedPredictiveModeling)
library(pROC)
library(MASS)
library(class)
#import data
data("Weekly")
dat = Weekly
head(dat)
```

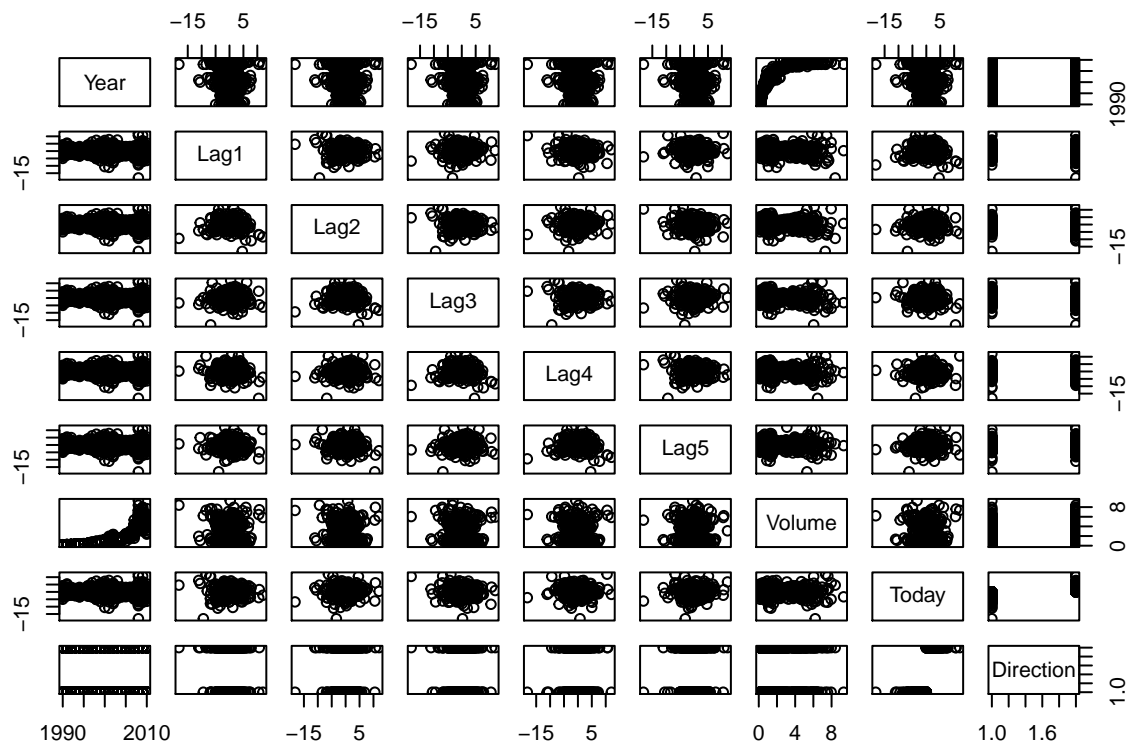
	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
## 1	1990	0.816	1.572	-3.936	-0.229	-3.484	0.1549760	-0.270	Down
## 2	1990	-0.270	0.816	1.572	-3.936	-0.229	0.1485740	-2.576	Down
## 3	1990	-2.576	-0.270	0.816	1.572	-3.936	0.1598375	3.514	Up
## 4	1990	3.514	-2.576	-0.270	0.816	1.572	0.1616300	0.712	Up
## 5	1990	0.712	3.514	-2.576	-0.270	0.816	0.1537280	1.178	Up
## 6	1990	1.178	0.712	3.514	-2.576	-0.270	0.1544440	-1.372	Down

(a) Produce some graphical summaries of the *Weekly* data.

```
transparentTheme(trans = .4)
featurePlot(x = dat[, 1:8],
            y = dat$Direction,
            scales = list(x=list(relation="free"),
                          y=list(relation="free")),
            plot = "density", pch = "|",
            auto.key = list(columns = 2))
```



```
pairs(dat)
```



(b) Use the full data set to perform a logistic regression with *Direction* as the response and the five Lag variables plus *Volume* as predictors. Do any of the predictors appear to be statistically significant? If so, which ones?

```
glm.fit <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,
               data=dat, family="binomial")
summary(glm.fit)
```

```
##
## Call:
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
##      Volume, family = "binomial", data = dat)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6949  -1.2565   0.9913   1.0849   1.4579
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.26686    0.08593   3.106  0.0019 **
## Lag1        -0.04127    0.02641  -1.563  0.1181
## Lag2         0.05844    0.02686   2.175  0.0296 *
## Lag3        -0.01606    0.02666  -0.602  0.5469
## Lag4        -0.02779    0.02646  -1.050  0.2937
## Lag5        -0.01447    0.02638  -0.549  0.5833
## Volume      -0.02274    0.03690  -0.616  0.5377
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1496.2  on 1088  degrees of freedom
## Residual deviance: 1486.4  on 1082  degrees of freedom
## AIC: 1500.4
##
## Number of Fisher Scoring iterations: 4
```

Lag2 appears to be statistically significant with p-value 0.0296, which is less than 0.05.

(c) Compute the confusion matrix and overall fraction of correct predictions. Briefly explain what the confusion matrix is telling you.

```
test.pred.prob <- predict(glm.fit, type = "response")
test.pred <- rep("Down", length(test.pred.prob))
test.pred[test.pred.prob>0.5] <- "Up"
confusionMatrix(data = as.factor(test.pred), reference = dat$Direction)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction Down  Up
##      Down    54  48
##      Up     430 557
##
```

```

##              Accuracy : 0.5611
##              95% CI : (0.531, 0.5908)
##      No Information Rate : 0.5556
##      P-Value [Acc > NIR] : 0.369
##
##              Kappa : 0.035
##
##      McNemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.11157
##              Specificity : 0.92066
##      Pos Pred Value : 0.52941
##      Neg Pred Value : 0.56434
##              Prevalence : 0.44444
##      Detection Rate : 0.04959
##      Detection Prevalence : 0.09366
##      Balanced Accuracy : 0.51612
##
##      'Positive' Class : Down
##

```


A confusion matrix is a table that is often used to describe the performance of a classification model (or “classifier”) on a set of test data for which the true values are known. This confusion matrix tells us that (1) the percentage of correct predictions on the training data is 56.11%, or say, the training error rate is 43.89%. (2) For weeks when the market goes up, the model is right 92.07% of the time; for weeks when the market goes down, the model is right 11.16% of the time.

(d) Plot the ROC curve using the predicted probability from logistic regression and report the AUC.

```

roc.glm <- roc(dat$Direction, test.pred.prob)
plot(roc.glm, legacy.axes = TRUE, print.auc = TRUE)
plot(smooth(roc.glm), col = 4, add = TRUE)

```

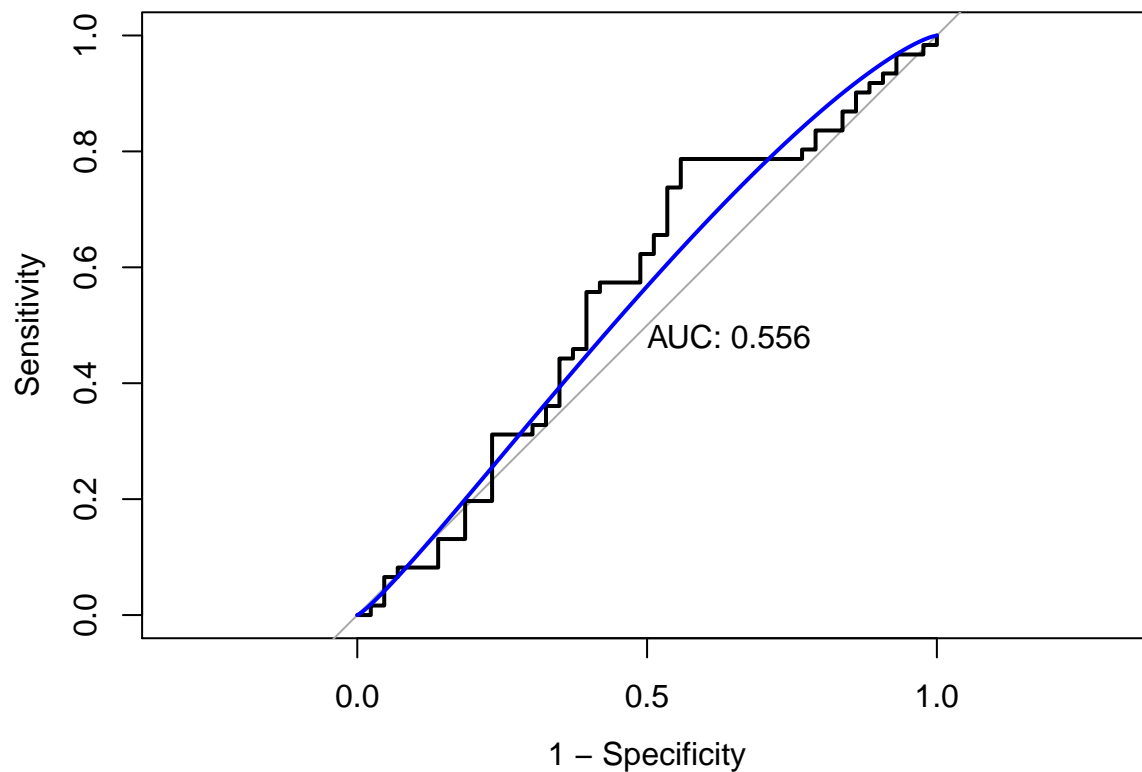


The AUC is 0.554.

(e) Now fit the logistic regression model using a training data period from 1990 to 2008, with *Lag1* and *Lag2* as the predictors. Plot the ROC curve using the held out data (that is, the data from 2009 and 2010) and report the AUC.

```
trainset = (dat$Year<=2008)
testset = dat[!trainset,]

glm.fit.d <- glm(Direction ~ Lag1 + Lag2, data=dat, subset=trainset, family="binomial")
glm.probs.d <- predict(glm.fit.d, type="response", newdata=testset)
roc.glm <- roc(testset$Direction, glm.probs.d)
plot(roc.glm, legacy.axes = TRUE, print.auc = TRUE)
plot(smooth(roc.glm), col = 4, add = TRUE)
```



The AUC is 0.556.


(f) Repeat (e) using LDA and QDA.

```
# LDA
lda.fit <- lda(Direction ~ Lag1 + Lag2, data=dat, subset=trainset)
lda.pred <- predict(lda.fit, newdata = testset)
head(lda.pred$posterior)
```

```
##          Down          Up
## 986 0.5602039 0.4397961
## 987 0.3079163 0.6920837
## 988 0.4458032 0.5541968
## 989 0.4785107 0.5214893
## 990 0.4657943 0.5342057
## 991 0.5262907 0.4737093
```

```
roc.lda <- roc(testset$Direction, lda.pred$posterior[,2],
               levels = c("Down", "Up"))
```

```
plot(roc.lda, legacy.axes = TRUE, print.auc = TRUE)
```




For LDA, the AUC is 0.557.

```
# QDA
qda.fit <- qda(Direction ~ Lag1 + Lag2, data=dat, subset=trainset)
qda.pred <- predict(qda.fit, newdata = testset)
head(qda.pred$posterior)
```

```
##           Down           Up
## 986 0.5436205 0.4563795
## 987 0.3528814 0.6471186
## 988 0.2227273 0.7772727
## 989 0.3483016 0.6516984
## 990 0.4598550 0.5401450
## 991 0.5119613 0.4880387
```

```
roc.qda <- roc(testset$Direction, qda.pred$posterior[,2],
               levels = c("Down", "Up"))
plot(roc.qda, legacy.axes = TRUE, print.auc = TRUE)
```



For QDA, the AUC is 0.529.

(g) Repeat (e) using KNN. Briefly discuss your results.

```
# choose the best K
train = dat %>%
  filter(Year<=2008)

test = dat %>%
  filter(Year>2008)

ctrl <- trainControl(method="repeatedcv",repeats = 3,classProbs=TRUE,summaryFunction = twoClassSummary)
knnFit <- train(Direction ~ Lag1 + Lag2, data = train,
  method = "knn",
  trControl = ctrl,
  preProcess = c("center","scale"),
  tuneLength = 20)
```

```
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was
## not in the result set. ROC will be used instead.
```

```
knnFit
```

```
## k-Nearest Neighbors
##
## 985 samples
## 2 predictor
## 2 classes: 'Down', 'Up'
##
```




```

## Pre-processing: centered (2), scaled (2)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 887, 887, 886, 886, 887, 887, ...
## Resampling results across tuning parameters:
##
##   k   ROC      Sens      Spec
##   5  0.5456473  0.4467340  0.6317508
##   7  0.5457662  0.4225253  0.6391807
##   9  0.5378838  0.3998653  0.6495735
##  11  0.5311326  0.3869697  0.6441639
##  13  0.5202741  0.3612963  0.6477890
##  15  0.5161328  0.3589899  0.6421437
##  17  0.5083680  0.3597811  0.6446128
##  19  0.5068425  0.3642256  0.6380022
##  21  0.5039217  0.3521717  0.6361279
##  23  0.5029667  0.3545455  0.6508754
##  25  0.5046668  0.3523064  0.6594388
##  27  0.5116593  0.3469697  0.6686195
##  29  0.5189606  0.3431145  0.6648822
##  31  0.5226016  0.3445791  0.6691582
##  33  0.5269912  0.3415993  0.6630415
##  35  0.5309649  0.3347811  0.6783838
##  37  0.5347202  0.3446465  0.6795511
##  39  0.5355586  0.3370707  0.6923906
##  41  0.5383194  0.3423906  0.6887430
##  43  0.5373038  0.3408249  0.6934792
##
## ROC was used to select the optimal model using the largest value.
## The final value used for the model was k = 7.

```

```
plot(knnFit)
```



```
# K=7 has the highest accuracy rate
```

```
knnPredict <- predict(knnFit,newdata = test , type="prob")
```

```
knnROC <- roc(test$Direction, knnPredict[, "Down"])
```

```
knnROC
```

```
##
```

```
## Call:
```

```
## roc.default(response = test$Direction, predictor = knnPredict[, "Down"])
```

```
##
```

```
## Data: knnPredict[, "Down"] in 43 controls (test$Direction Down) > 61 cases (test$Direction Up).
```

```
## Area under the curve: 0.5448
```

```
plot(knnROC, type="S", print.auc = TRUE)
```

